

Practical No - 1

Aim:List of Commands(mkdir, touchz, copy from local/put,copy to local/get move from local,cp,rmr,du,dus,stat)

start-all.sh

```
hadoop@lab263-B250M-D2V:~$ start-all.sh
WARNING: Attempting to start all Apache Hadoop daemons as hadoop in 10 seconds.
WARNING: This is not a recommended production deployment configuration.
WARNING: Use CTRL-C to abort.
Starting namenodes on [localhost]
Starting datanodes
Starting secondary namenodes [lab263-B250M-D2V]
Starting resourcemanager
Starting nodemanagers
hadoop@lab263-B250M-D2V:~$
```

1)hadoop fs

hadoop fs -ls

```
hadoop@lab263-B250M-D2V:~$ hadoop fs -ls
Found 1 items
-rw-r--r--  1 hadoop supergroup          0 2023-11-09 15:09 Demo.txt
hadoop@lab263-B250M-D2V:~$
```

2)touchz: It creates an empty file.

hadoop fs -mkdir -p /user/hadoop/

hadoop fs -touchz /user/hadoop/Demo.txt

```
hadoop@lab263-B250M-D2V:~$ hadoop fs -mkdir -p /user/hadoop/
hadoop@lab263-B250M-D2V:~$ hadoop fs -touchz /user/hadoop/Demo.txt
hadoop@lab263-B250M-D2V:~$ hadoop fs -ls
Found 1 items
-rw-r--r--  1 hadoop supergroup          0 2023-11-09 15:09 Demo.txt
hadoop@lab263-B250M-D2V:~$
```

3)copyFromLocal (or) put:

hadoop fs -copyFromLocal test.txt Demo.txt

```
hadoop@lab263-B250M-D2V:~$ hadoop fs -copyFromLocal test.txt Demo.txt
copyFromLocal: `Demo.txt': File exists
hadoop@lab263-B250M-D2V:~$
```

4) test

hadoop fs -test -d Demo

echo s7

```
hadoop@lab263-B250M-D2V:~$ hadoop fs -test -d Demo
hadoop@lab263-B250M-D2V:~$
hadoop@lab263-B250M-D2V:~$ echo s7
s7
hadoop@lab263-B250M-D2V:~$
```


```
hadoop@lab263-B250M-D2V:~$ nano test1.txt
hadoop@lab263-B250M-D2V:~$ nano test2.txt
hadoop@lab263-B250M-D2V:~$ hadoop fs -touchz Demo.txt
hadoop@lab263-B250M-D2V:~$ hadoop fs -appendToFile test1.txt test2.txt Demo.txt
hadoop@lab263-B250M-D2V:~$ hadoop fs -cat Demo.txt
Texttttttt 1
Texttttttt 2
hadoop@lab263-B250M-D2V:~$
```

5)mkdir

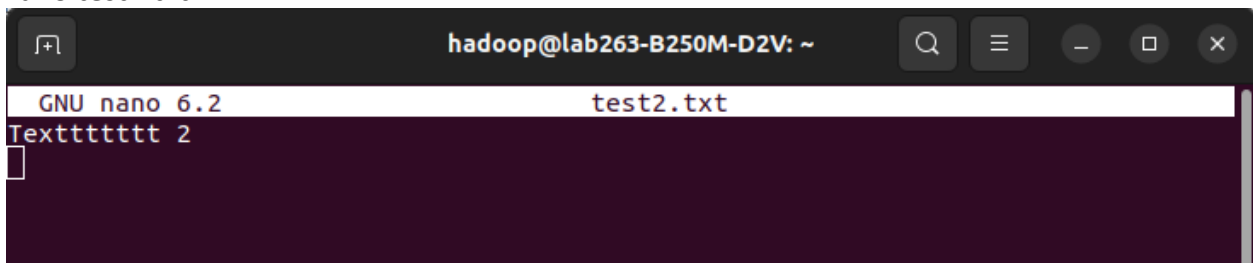
hadoop fs -mkdir Demodirectory

```
hadoop@lab263-B250M-D2V:~$ hadoop fs -mkdir Demodirectory
hadoop@lab263-B250M-D2V:~$
hadoop@lab263-B250M-D2V:~$
```

nano test1.txt

A screenshot of a terminal window with the nano editor open. The title bar shows 'hadoop@lab263-B250M-D2V: ~'. The editor window shows 'GNU nano 6.2' and 'test1.txt *'. The content of the file is 'Texttttttt 1' followed by a cursor.

nano test2.txt

A screenshot of a terminal window with the nano editor open. The title bar shows 'hadoop@lab263-B250M-D2V: ~'. The editor window shows 'GNU nano 6.2' and 'test2.txt'. The content of the file is 'Texttttttt 2' followed by a cursor.

6)appendToFile

nano test1.txt

nano test2.txt

hadoop fs -touchz Demo.txt

hadoop fs -appendToFile test1.txt test2.txt Demo.txt

hadoop fs -cat Demo.txt

```
hadoop@lab263-B250M-D2V:~$ nano test1.txt
hadoop@lab263-B250M-D2V:~$ nano test2.txt
hadoop@lab263-B250M-D2V:~$ hadoop fs -touchz Demo.txt
hadoop@lab263-B250M-D2V:~$ hadoop fs -appendToFile test1.txt test2.txt Demo.txt
hadoop@lab263-B250M-D2V:~$ hadoop fs -cat Demo.txt
Texttttttt 1
Texttttttt 2
hadoop@lab263-B250M-D2V:~$
```

7)usage

hadoop fs -usage test

```
hadoop@lab263-B250M-D2V:~$ hadoop fs -usage test
Usage: hadoop fs [generic options] -test -[defswrz] <path>
hadoop@lab263-B250M-D2V:~$
```

8)Count

hadoop fs -count -v /

```
hadoop@lab263-B250M-D2V:~$ hadoop fs -count -v /
DIR_COUNT  FILE_COUNT  CONTENT_SIZE  PATHNAME
      19         50      9772997  /
hadoop@lab263-B250M-D2V:~$
```

9)find

hadoop fs -find / -name Demodirectory

```
hadoop@lab263-B250M-D2V:~$ hadoop fs -count -v /
DIR_COUNT  FILE_COUNT  CONTENT_SIZE  PATHNAME
      19         50      9772997  /
hadoop@lab263-B250M-D2V:~$
```

10)help

hadoop fs -help

count

```
hadoop@lab263-B250M-D2V:~$ hadoop fs -help count
-count [-q] [-h] [-v] [-t [<storage type>]] [-u] [-x] [-e] [-s] <path> ... :
  Count the number of directories, files and bytes under the paths
  that match the specified file pattern.  The output columns are:
  DIR_COUNT FILE_COUNT CONTENT_SIZE PATHNAME
  or, with the -q option:
  QUOTA REM_QUOTA SPACE_QUOTA REM_SPACE_QUOTA
    DIR_COUNT FILE_COUNT CONTENT_SIZE PATHNAME
  The -h option shows file sizes in human readable format.
  The -v option displays a header line.
```

Practical No - 2

Aim: Write a Program in Map Reduce for WordCount operation.

WordCountDriver.java

```
import org.apache.hadoop.fs.Path; import
org.apache.hadoop.io.Text; import
org.apache.hadoop.io.IntWritable;

import org.apache.hadoop.mapreduce.Job;

import org.apache.hadoop.mapreduce.lib.input.FileInputFormat; import
org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;import
org.apache.hadoop.mapreduce.Reducer;

import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.conf.Configuration;
public class WordCountDriver {

    public static void main(String[] args) throws Exception{Job
        j1=Job.getInstance(new Configuration());
        j1.setJarByClass(WordCountDriver.class);
        j1.setJobName("Average Word Count");

        FileInputFormat.addInputPath(j1,new Path(args[0]));
        FileOutputFormat.setOutputPath(j1, new Path(args[1]));

        j1.setMapperClass(WordCountMapper.class);
        j1.setReducerClass(WordCountReducer.class);
        j1.setOutputKeyClass(Text.class); j1.setOutputValueClass(IntWritable.class);
        System.exit(j1.waitForCompletion(true)? 0:1);
    }
}
```

WordCountMapper.java

```
import java.io.IOException; import
java.util.StringTokenizer; import
org.apache.hadoop.io.Text;

import org.apache.hadoop.io.IntWritable; import
org.apache.hadoop.io.LongWritable; import
org.apache.hadoop.mapreduce.Mapper; import
org.apache.hadoop.mapreduce.Reducer;

import org.apache.hadoop.mapreduce.Reducer.Context;

public class WordCountMapper extends
Mapper<LongWritable,Text,Text,IntWritable> {

    private final static IntWritable one=new IntWritable(1)

private Text word=new Text();

    public void map(LongWritable key, Text value, Context context) throws
IOException,InterruptedException {

        String line=value.toString();

        StringTokenizer ltr=new StringTokenizer(line);
        while(ltr.hasMoreTokens()){

            word.set(ltr.nextToken());
            context.write(word ,one);

        }

    }

}
```

WordCountReducer.java

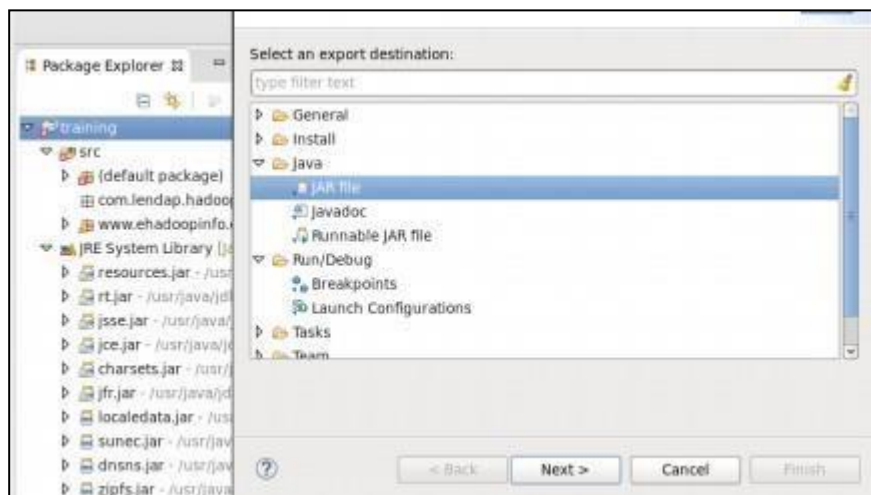
```
import java.io.IOException; import
org.apache.hadoop.io.Text;

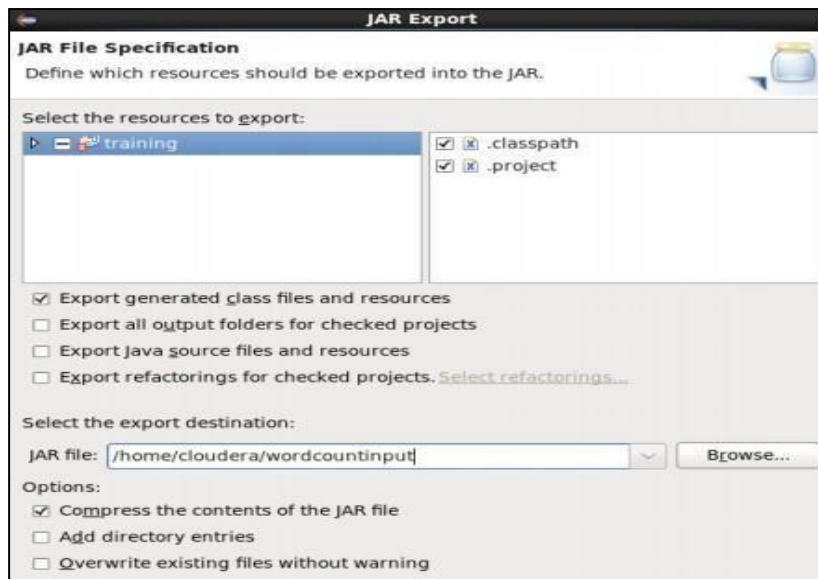
import org.apache.hadoop.io.IntWritable; import
org.apache.hadoop.mapreduce.Reducer; public class
WordCountReducer extends
Reducer<Text,IntWritable,Text,IntWritable> {

    public void reduce(Text key,Iterable<IntWritable> values,Context
context) throws IOException,InterruptedException{

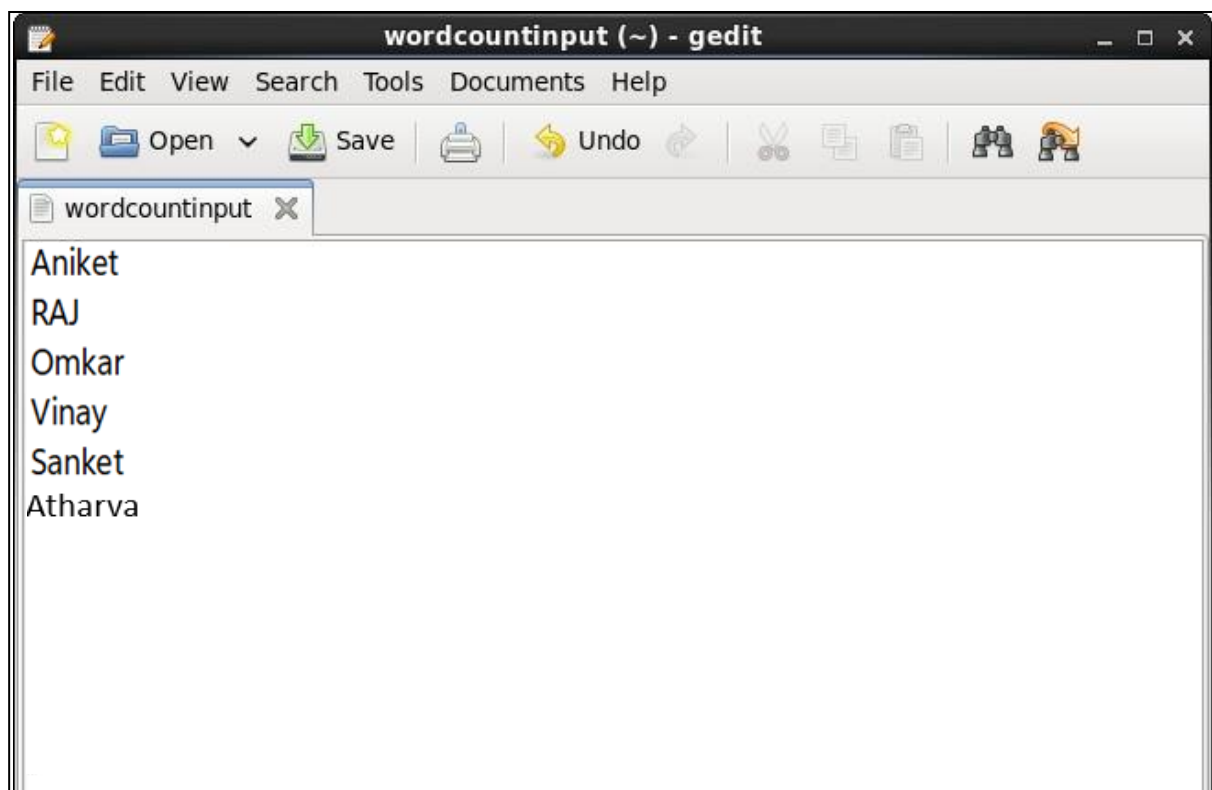
        int sum=0;
        for(IntWritable value:values)
        {
            sum+=value.get();
        }
        context.write(key, new IntWritable(sum));
    }
}
```

Export .jar file. Right click on training and select export.





Follow below HDFS commands



```
[cloudera@quickstart ~]$ sudo -u hdfs hadoop fs -copyFromLocal wordcountinput /wordcountinputhdfs
[cloudera@quickstart ~]$
[cloudera@quickstart ~]$
[cloudera@quickstart ~]$ hadoop fs cat /wordcountinputhdfs
cat: Unknown command
Did you mean -cat? This command begins with a dash.
[cloudera@quickstart ~]$ hadoop fs -cat /wordcountinputhdfs
```

Aniket

RAJ

Omkar

Vinay

Sanket

Atharva

Executing the jar file using hadoop command:

```
cloudera@quickstart:~$ sudo -u hdfs hadoop jar wordcountinput.jar WordCountDriver /wordcountinputhdfs /wordcountoutputdir3
22/11/13 03:27:57 INFO client.RMPProxy: Connecting to ResourceManager at quickstart.cloudera/10.0.2.15:8032
22/11/13 03:27:58 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface
execute your application with ToolRunner to remedy this.
22/11/13 03:27:58 INFO input.FileInputFormat: Total input paths to process : 1
22/11/13 03:27:58 INFO mapreduce.JobSubmitter: number of splits:1
22/11/13 03:27:58 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1668337644684_0005
22/11/13 03:27:58 INFO impl.YarnClientImpl: Submitted application application_1668337644684_0005
22/11/13 03:27:58 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8088/proxy/application_1668337644684_0005
22/11/13 03:27:58 INFO mapreduce.Job: Running job: job_1668337644684_0005
22/11/13 03:28:06 INFO mapreduce.Job: Job job_1668337644684_0005 running in uber mode : false
22/11/13 03:28:06 INFO mapreduce.Job:  map 0% reduce 0%
22/11/13 03:28:12 INFO mapreduce.Job:  map 100% reduce 0%
22/11/13 03:28:18 INFO mapreduce.Job:  map 100% reduce 100%
22/11/13 03:28:19 INFO mapreduce.Job: Job job_1668337644684_0005 completed successfully
22/11/13 03:28:19 INFO mapreduce.Job: Counters: 49
  File System Counters
    FILE: Number of bytes read=146
    FILE: Number of bytes written=256611
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=224
    HDFS: Number of bytes written=88
    HDFS: Number of read operations=6
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
```


Aim: write a Program in Map Reduce for Matrix Multiplication.

MatrixMultiplication.java

```
import java.io.DataInput;
import java.io.DataOutput;
import java.io.IOException;
import java.util.ArrayList;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;

import org.apache.hadoop.io.DoubleWritable;
import org.apache.hadoop.io.IntWritable; import
org.apache.hadoop.io.Writable;

import org.apache.hadoop.io.WritableComparable;
import org.apache.hadoop.io.Text;

import org.apache.hadoop.mapreduce.Job; import
org.apache.hadoop.mapreduce.Mapper; import
org.apache.hadoop.mapreduce.Reducer; import
org.apache.hadoop.mapreduce.lib.input.*;

import org.apache.hadoop.mapreduce.lib.output.*;
import org.apache.hadoop.util.ReflectionUtils;

class Element implements Writable {int
    tag;

    int index;
    double value;
    Element() {

        tag = 0;
        index = 0;
        value = 0.0;
    }
    Element(int tag, int index, double value) {
        this.tag = tag;
        this.index = index;
        this.value = value;
    }
}
```

```
@Override
public void readFields(DataInput input) throws IOException { tag =
    input.readInt();
    index = input.readInt(); value
    = input.readDouble();
}
@Override
public void write(DataOutput output) throws IOException {
    output.writeInt(tag);
output.writeInt(index);
output.writeDouble(value);
}
}
class Pair implements WritableComparable<Pair> {int i;
    int j;
    Pair() {
        i = 0;
        j = 0;
    }
    Pair(int i, int j) {
        this.i = i;
        this.j = j;
    }
    @Override
    public void readFields(DataInput input) throws IOException { i =
        input.readInt();
        j = input.readInt();
    }
    @Override
    public void write(DataOutput output) throws IOException {

        output.writeInt(i);
        output.writeInt(j);
    }
}
```

```
@Override
public int compareTo(Pair compare) {if (i
    > compare.i) {
        return 1;
    }
    else if (i < compare.i) {
        return -1;
    }
    else {
        if (j > compare.j) {
            return 1;
        }
        else if (j < compare.j) {
            return -1;
        }
    }
    return 0;
}

public String toString() { return
    i + " " + j + " ";
}
}

public class MatrixMultiply {
    public static class MatrixMapperM extends Mapper<Object,
        Text, IntWritable, Element> {

        @Override
        public void map(Object key, Text value, Context context)
            throws IOException, InterruptedException {

            String readLine = value.toString();
            String[] tokens = readLine.split(","); int
            index = Integer.parseInt(tokens[0]);

            double elementVal = Double.parseDouble(tokens[2]);
            Element e = new Element(0, index, elementVal);

            IntWritable keyval = new IntWritable(Integer.parseInt(tokens[1]));
            context.write(keyval, e);
        }
    }
}
```

```
public static class MatrixMapperN extends Mapper<Object,
    Text, IntWritable, Element> {

    @Override
    public void map(Object key, Text value, Context context)
        throws IOException, InterruptedException {

        String readLine = value.toString();
        String[] tokens = readLine.split(","); int
        index = Integer.parseInt(tokens[1]);

        double elementVal = Double.parseDouble(tokens[2]);
        Element e = new Element(1, index, elementVal);

        IntWritable keyval = new IntWritable(Integer.parseInt(tokens[0]));
        context.write(keyval, e);
    }
}

public static class ReducerMN extends
    Reducer<IntWritable, Element, Pair, DoubleWritable> {

    @Override
    public void reduce(IntWritable key, Iterable<Element> values,
        Context context) throws IOException, InterruptedException {
        ArrayList<Element> M = new ArrayList<Element>();
        ArrayList<Element> N = new ArrayList<Element>(); Configuration
        conf = context.getConfiguration();
        for (Element element : values) {
            Element temp = ReflectionUtils.newInstance(Element.class, conf);
            ReflectionUtils.copy(conf, element, temp);
            if (temp.tag == 0)
            {
                M.add(temp);
            }
            else if (temp.tag == 1)
            {
                N.add(temp);
            }
        }
    }
}
```

```
        for (int i = 0; i < M.size(); i++) {
            for (int j = 0; j < N.size(); j++) {
                Pair p = new Pair(M.get(i).index, N.get(j).index);
                double mul = M.get(i).value * N.get(j).value;

                context.write(p, new DoubleWritable(mul));
            }
        }
    }

}

public static class MapMN extends
    Mapper<Object, Text, Pair, DoubleWritable> {
    @Override
    public void map(Object key, Text value, Context context)
        throws IOException, InterruptedException {

        String readLine = value.toString(); String[]
        pairValue = readLine.split(" ");

        Pair p = new Pair(Integer.parseInt(pairValue[0]),
            Integer.parseInt(pairValue[1]));

        DoubleWritable val = new DoubleWritable(
            Double.parseDouble(pairValue[2]));

        context.write(p, val);
    }
}

public static class ReduceMN extends
    Reducer<Pair, DoubleWritable, Pair, DoubleWritable> {
    @Override
    public void reduce(Pair key, Iterable<DoubleWritable> values, Context
        context) throws IOException, InterruptedException {

        double sum = 0.0;
        for (DoubleWritable value : values) {
            sum += value.get();
        }

        context.write(key, new DoubleWritable(sum));
    }
}
```

```

public static void main(String[] args) throws Exception {
    Path MPath = new Path("/expt4/input/M"); Path NPath = new Path("/expt4/input/N");
    Path intermediatePath = new Path("/expt4/interim"); Path
    outputPath = new Path("/expt4/output");

    Job job1 = Job.getInstance();
    job1.setJobName("Map Intermediate");
    job1.setJarByClass(MatrixMultiply.class);

    MultipleInputs.addInputPath(job1, MPath, TextInputFormat.class,
    MatrixMapperM.class);

    MultipleInputs.addInputPath(job1, NPath, TextInputFormat.class,
    MatrixMapperN.class);

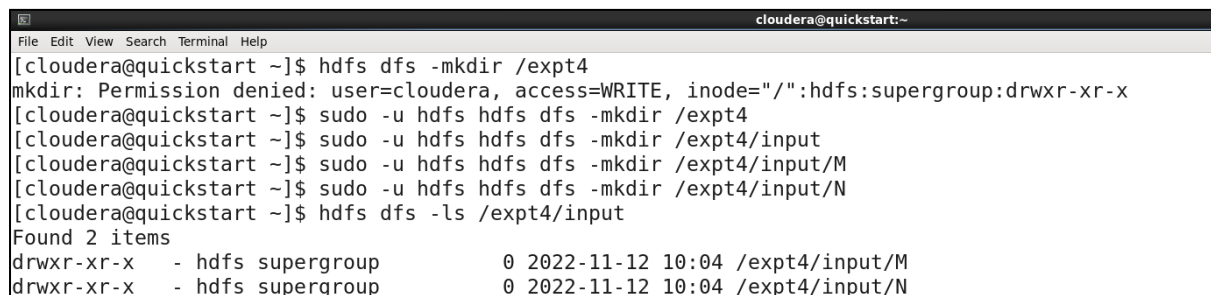
    job1.setReducerClass(ReducerMN.class);
    job1.setMapOutputKeyClass(IntWritable.class);
    job1.setMapOutputValueClass(Element.class); job1.setOutputKeyClass(Pair.class);
    job1.setOutputValueClass(DoubleWritable.class);
    job1.setOutputFormatClass(TextOutputFormat.class);
    FileOutputFormat.setOutputPath(job1, intermediatePath); job1.waitForCompletion(true);

    Job job2 = Job.getInstance(); job2.setJobName("Map
    Final Output");
    job2.setJarByClass(MatrixMultiply.class);
    job2.setMapperClass(MapMN.class);
    job2.setReducerClass(ReducerMN.class);
    job2.setOutputKeyClass(Pair.class);
    job2.setOutputValueClass(DoubleWritable.class);
    job2.setInputFormatClass(TextInputFormat.class);

    job2.setOutputFormatClass(TextOutputFormat.class);
    FileInputFormat.addInputPath(job2, intermediatePath);
    FileOutputFormat.setOutputPath(job2, outputPath);
    job2.waitForCompletion(true);
}
}

```

Prerequisites create the input directories to store the input matrices M and N

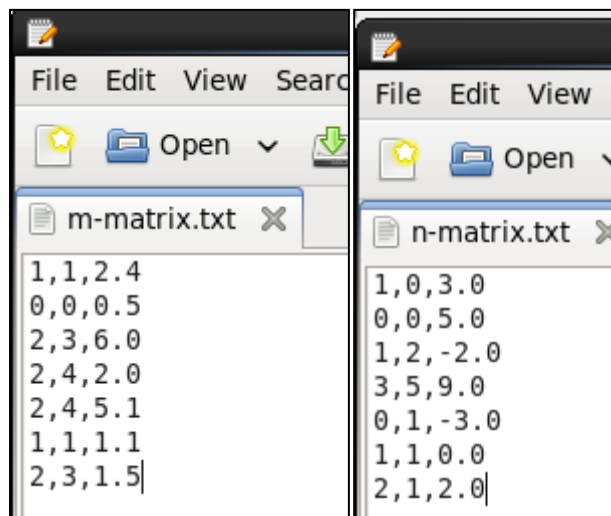


```

cloudera@quickstart:~$ hdfs dfs -mkdir /expt4
mkdir: Permission denied: user=cloudera, access=WRITE, inode="/":hdfs:supergroup:drwxr-xr-x
cloudera@quickstart:~$ sudo -u hdfs hdfs dfs -mkdir /expt4
cloudera@quickstart:~$ sudo -u hdfs hdfs dfs -mkdir /expt4/input
cloudera@quickstart:~$ sudo -u hdfs hdfs dfs -mkdir /expt4/input/M
cloudera@quickstart:~$ sudo -u hdfs hdfs dfs -mkdir /expt4/input/N
cloudera@quickstart:~$ hdfs dfs -ls /expt4/input
Found 2 items
drwxr-xr-x - hdfs supergroup          0 2022-11-12 10:04 /expt4/input/M
drwxr-xr-x - hdfs supergroup          0 2022-11-12 10:04 /expt4/input/N

```

The figure below shows the matrix data used for this implementation:



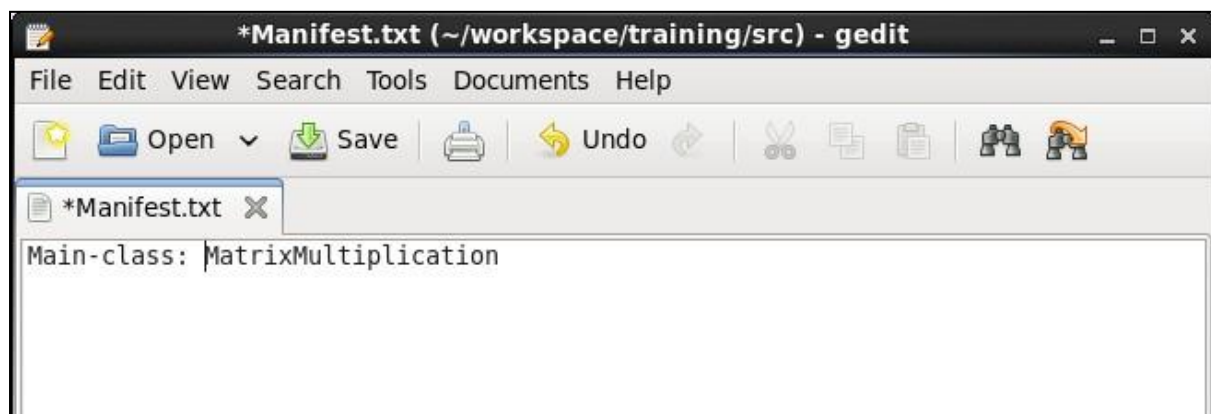
Copy the matrix data from the local system to HDFS

```
[cloudera@quickstart ~]$ sudo -u hdfs hdfs dfs -copyFromLocal m-matrix.txt /expt4/input/M
[cloudera@quickstart ~]$
[cloudera@quickstart ~]$ sudo -u hdfs hdfs dfs -copyFromLocal n-matrix.txt /expt4/input/N
[cloudera@quickstart ~]$
[cloudera@quickstart ~]$ hdfs dfs -ls /expt4/input/M
Found 1 items
-rw-r--r-- 1 hdfs supergroup 56 2022-11-12 10:12 /expt4/input/M/m-matrix.txt
[cloudera@quickstart ~]$ hdfs dfs -ls /expt4/input/N
Found 1 items
-rw-r--r-- 1 hdfs supergroup 58 2022-11-12 10:12 /expt4/input/N/n-matrix.txt
[cloudera@quickstart ~]$
[cloudera@quickstart ~]$
```

Compile the code to create the classes

```
cloudera@quickstart:~/workspace/training/src
File Edit View Search Terminal Help
[cloudera@quickstart src]$ javac MatrixMultiplication.java -cp $(hadoop classpath)
[cloudera@quickstart src]$
[cloudera@quickstart src]$
[cloudera@quickstart src]$ gedit Manifest.txt
[cloudera@quickstart src]$
[cloudera@quickstart src]$
```

To indicate the main class file, create a Manifest file to point out to the main driver class.



Compile and create the Jar file required to run the MapReduce Task

```
cloudera@quickstart:~/workspace/training/src
File Edit View Search Terminal Help
[cloudera@quickstart src]$ jar -cfm MatrixMultiplication.jar Manifest.txt *.class
[cloudera@quickstart src]$
```

Run the jar file on the Hadoop ecosystem to trigger all the MapReduce classes.

```
cloudera@quickstart:~/workspace/training/src
File Edit View Search Terminal Help
[cloudera@quickstart src]$ sudo -u hdfs hadoop jar MatrixMultiplication.jar
22/11/12 10:29:23 INFO client.RMPProxy: Connecting to ResourceManager at quickstart.cloudera/127.0.0.1:8032
22/11/12 10:29:24 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
22/11/12 10:29:24 INFO input.FileInputFormat: Total input paths to process : 1
22/11/12 10:29:24 INFO input.FileInputFormat: Total input paths to process : 1
22/11/12 10:29:24 INFO mapreduce.JobSubmitter: number of splits:2
22/11/12 10:29:24 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1668275009795_0006
22/11/12 10:29:24 INFO impl.YarnClientImpl: Submitted application application_1668275009795_0006
22/11/12 10:29:24 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8088/proxy/application_1668275009795_0006/
22/11/12 10:29:24 INFO mapreduce.Job: Running job: job_1668275009795_0006
22/11/12 10:29:31 INFO mapreduce.Job: Job job_1668275009795_0006 running in uber mode : false
22/11/12 10:29:31 INFO mapreduce.Job: map 0% reduce 0%
22/11/12 10:29:41 INFO mapreduce.Job: map 100% reduce 0%
22/11/12 10:29:48 INFO mapreduce.Job: map 100% reduce 100%
22/11/12 10:29:48 INFO mapreduce.Job: Job job_1668275009795_0006 completed successfully
22/11/12 10:29:48 INFO mapreduce.Job: Counters: 49
File System Counters
  FILE: Number of bytes read=175
  FILE: Number of bytes written=387639
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=648
  HDFS: Number of bytes written=124
  HDFS: Number of read operations=9
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=2
```

Output of the file

```
cloudera@quickstart:~/workspace/training
File Edit View Search Terminal Help
[cloudera@quickstart training]$ hdfs dfs -cat /expt4/interim/part-r-00000
0 1 -1.5
0 0 2.5
1 1 0.0
1 2 -2.2
1 0 3.3000000000000003
1 1 0.0
1 2 -4.8
1 0 7.199999999999999
2 5 13.5
2 5 54.0
[cloudera@quickstart training]$ hdfs dfs -cat /expt4/output/part-r-00000
0 0 2.5
0 1 -1.5
1 0 10.5
1 1 0.0
1 2 -7.0
2 5 67.5
[cloudera@quickstart training]$
```


Practical No - 3

Aim: Query the Sample Database using MongoDB querring commands.

1. create a database named salesdb and switch to salesdb.
2. We have two collections namely sales and sales_profile. Let us populate them. First run a query to insert data in sales_profile.

```
test> use salesdb;
switched to db salesdb
salesdb> db.sales_profile.insertMany([
... { id:1,invoice_no:"750-67-8428", card:"Member", gender:"Female"} ,
... { id:2,invoice_no:"226-31-3081", card:"Normal", gender:"Female"} ,
... { id:3,invoice_no:"631-41-3108", card:"Normal", gender:"Male"} ,
... { id:4,invoice_no:"123-19-1176", card:"Member", gender:"Male"} ,
... { id:5,invoice_no:"373-73-7910", card:"Normal", gender:"Male"} ])
{
  acknowledged: true,
  insertedIds: {
    '0': ObjectId("633896ee748becd9a0a5da6c"),
    '1': ObjectId("633896ee748becd9a0a5da6d"),
    '2': ObjectId("633896ee748becd9a0a5da6e"),
    '3': ObjectId("633896ee748becd9a0a5da6f"),
    '4': ObjectId("633896ee748becd9a0a5da70")
  }
}
```

Then we run a query to import sales_profile.csv in sales_profile collection.

Query :mongoimport --db=salesdb --collection=sales --type=csv --headerline --file=sales_profile.csv

```
Microsoft Windows [Version 10.0.19044.2130]
(c) Microsoft Corporation. All rights reserved.

C:\Users\Lab1-89\Downloads\mongodb-database-tools-windows-x86_64-100.6.1\bin>mongoimport --db=salesdb --collection=sales -
--type=csv --headerline --file=sales_profile.csv
2022-11-14T11:52:07.541+0530    connected to: mongodb://localhost/
2022-11-14T11:52:07.892+0530    2823 document(s) imported successfully. 0 document(s) failed to import.

C:\Users\Lab1-89\Downloads\mongodb-database-tools-windows-x86_64-100.6.1\bin>
```

3. Find data query

Query : db.sales.find()

```

mongosh mongodb://localhost:27017/?directConnection=true&serverSelectionTimeoutMS=2000
salesdb> db.sales.find()
[
  {
    _id: ObjectId("6371de8f8976cbf6d0144748"),
    ORDERNUMBER: 10107,
    QUANTITYORDERED: 30,
    PRICEEACH: 95.7,
    ORDERLINENUMBER: 2,
    SALES: 2871,
    ORDERDATE: '2/24/2003 0:00',
    STATUS: 'Shipped',
    QTR_ID: 1,
    MONTH_ID: 2,
    YEAR_ID: 2003,
    PRODUCTLINE: 'Motorcycles',
    MSRP: 95,
    PRODUCTCODE: 'S10_1678',
    CUSTOMERNAME: 'Land of Toys Inc.',
    PHONE: 2125557818,
    ADDRESSLINE1: '897 Long Airport Avenue',
    ADDRESSLINE2: '',
    CITY: 'NYC',
    STATE: 'NY',
    POSTALCODE: 10022,
    COUNTRY: 'USA',
    TERRITORY: 'NA',
    CONTACTLASTNAME: 'Yu',
    CONTACTFIRSTNAME: 'Kwai',
    DEALSIZE: 'Small'
  }
]

```

db.sales.find({ CITY : "Salzburg" , STATUS : "Shipped" }).limit(5)

```

mongosh mongodb://localhost:27017/?directConnection=true&serverSelectionTimeoutMS=2000
salesdb> db.sales.find( { CITY : "Salzburg" , STATUS : "Shipped" }).limit(5)
[
  {
    _id: ObjectId("6371de8f8976cbf6d014475b"),
    ORDERNUMBER: 10341,
    QUANTITYORDERED: 41,
    PRICEEACH: 100,
    ORDERLINENUMBER: 9,
    SALES: 7737.93,
    ORDERDATE: '11/24/2004 0:00',
    STATUS: 'Shipped',
    QTR_ID: 4,
    MONTH_ID: 11,
    YEAR_ID: 2004,
    PRODUCTLINE: 'Motorcycles',
    MSRP: 95,
    PRODUCTCODE: 'S10_1678',
    CUSTOMERNAME: 'Salzburg Collectables',
    PHONE: '6562-9555',
    ADDRESSLINE1: 'Geislweg 14',
    ADDRESSLINE2: '',
    CITY: 'Salzburg',
    STATE: '',
    POSTALCODE: 5020,
    COUNTRY: 'Austria',
    TERRITORY: 'EMEA',
    CONTACTLASTNAME: 'Pipps',
    CONTACTFIRSTNAME: 'Georg',
    DEALSIZE: 'Large'
  }
]

```

4. Find data using greater than condition

db.sales.find({ QUANTITYORDERED: { \$gt:33 }, PRICEEACH: { \$gt:33 } }, { CITY : "Salzburg" , STATUS : "Shipped" }).limit(5)

```

mongosh mongodb://localhost:27017/?directConnection=true&serverSelectionTimeoutMS=2000
salesdb> db.sales.find( { QUANTITYORDERED: { $gt:33 }, PRICEEACH: { $gt:33 } }, { CITY : "Salzburg" , STATUS : "Shipped" }).limit(5)
[
  {
    _id: ObjectId("6371de8f8976cbf6d0144749"),
    STATUS: 'Shipped',
    CITY: 'Reims'
  },
  {
    _id: ObjectId("6371de8f8976cbf6d014474a"),
    STATUS: 'Shipped',
    CITY: 'Paris'
  },
  {
    _id: ObjectId("6371de8f8976cbf6d014474b"),
    STATUS: 'Shipped',
    CITY: 'Pasadena'
  }
]

```

5. Update document

```
db.sales.updateMany( { CITY : "Salzburg"}, { $set :{ CITY:"MUMBAI" } })
```

```
salesdb> db.sales.updateMany( { CITY : "Salzburg"},  
... { $set :{ CITY:"MUMBAI" } })  
{  
  acknowledged: true,  
  insertedId: null,  
  matchedCount: 40,  
  modifiedCount: 40,  
  upsertedCount: 0  
}  
salesdb> _
```

6. Delete document db.sales.deleteOne({CITY : "MUMBAI" })

```
salesdb> db.sales.deleteOne({CITY: "MUMBAI"})  
{ acknowledged: true, deletedCount: 1 }  
salesdb>
```

7. Aggregate function

```
db.sales.aggregate([ { $lookup: { from: "sales_profile", localField: "ID", foreignField: "ID", as:  
"QUNTITTY" } } ])
```

```
mongosh mongodb://localhost:27017/?directConnection=true&serverSelectionTimeoutMS=2000  
salesdb> db.sales.aggregate([ { $lookup: { from: "sales_profile", localField: "ID", foreignField: "ID", as: "QUNTITTY" } } ])  
[  
  {  
    _id: ObjectId("6371de8f8976cbf6d0144748"),  
    ORDERNUMBER: 10107,  
    QUANTITYORDERED: 30,  
    PRICEEACH: 95.7,  
    ORDERLINENUMBER: 2,  
    SALES: 2871,  
    ORDERDATE: '2/24/2003 0:00',  
    STATUS: 'Shipped',  
    QTR_ID: 1,  
    MONTH_ID: 2,  
    YEAR_ID: 2003,  
    PRODUCTLINE: 'Motorcycles',  
    MSRP: 95,  
    PRODUCTCODE: 'S10_1678',  
    CUSTOMERNAME: 'Land of Toys Inc.',  
    PHONE: 2125557818,  
    ADDRESSLINE1: '897 Long Airport Avenue',  
    ADDRESSLINE2: '',  
    CITY: 'NYC',  
    STATE: 'NY',  
    POSTALCODE: 10022,  
    COUNTRY: 'USA',  
    TERRITORY: 'NA',  
    CONTACTLASTNAME: 'Yu',  
    CONTACTFIRSTNAME: 'Kwai',  
    DEALSIZE: 'Small',  
    QUNTITTY: []  
  }  
]
```

8. Extract the salesdb database in json

```
mongoexport --collection=sales --db=salesdb --out=sales.json
```

```
C:\Users\Lab1-89\Downloads\mongodb-database-tools-windows-x86_64-100.6.1\bin>mongoexport --collection=sales --db=salesdb --out=sales.json  
2022-11-14T12:59:26.100+0530 connected to: mongodb://localhost/  
2022-11-14T12:59:26.282+0530 exported 2822 records  
C:\Users\Lab1-89\Downloads\mongodb-database-tools-windows-x86_64-100.6.1\bin>_
```

Practical No : 4

Aim: Basic HIVE Commands

1. Log in into Hive.

```
cloudera-quickstart-vm-5.13.0-0-virtualbox [Running] - Oracle VM VirtualBox
File Machine View Input Devices Help
Applications Places System Wed Oct 16, 10:24:48 AM
Browse and run installed applications cloudera@quickstart:~
File Edit View Search Terminal Help
[cloudera@quickstart ~]$ sudo -uhive hive

Logging initialized using configuration in file:/etc/hive/conf.dist/hive-log4j.properties
WARNING: Hive CLI is deprecated and migration to Beeline is recommended.
```

2. Show available databases.

```
hive> show databases;
OK
default
demo
demo1
demo2
demo3
Time taken: 0.258 seconds, Fetched: 5 row(s)
```

3. Select a database and create a table name student.

```
hive> use demo;
OK
Time taken: 0.017 seconds
hive> create table student(name string, rollno int, dept string)
> row format delimited
> fields terminated by ',';
OK
Time taken: 0.187 seconds
hive>
```

4. Import data from a file and display the table with the data.

```
hive> load data local inpath '/home/cloudera/Desktop/student' into table student;
Loading data to table demo.student
Table demo.student stats: [numFiles=1, totalSize=70]
OK
Time taken: 0.443 seconds
hive> select * from student;
OK
Prajyot 53 MCA
Tushar 14 MCA
Aniket 42 MCA
Vinay 10 MCA
Sanket 55 MCA
Time taken: 0.146 seconds, Fetched: 5 row(s)
hive>
```

5. Show available tables.

```
hive> show tables;
OK
employee
student
Time taken: 0.017 seconds, Fetched: 2 row(s)
```

6. Describe the table.

```
hive> describe student;
OK
name                string
rollno              int
dept                string
Time taken: 0.035 seconds, Fetched: 3 row(s)
hive> █
```

USING HUE EDITOR

7. Create a database and a table.

Input:

CREATE DATABASE MCA;

USE MCA;

CREATE TABLE STUDENT(ID INT, NAME VARCHAR(10), FIELD VARCHAR(10));

INSERT INTO STUDENT VALUES(1, 'Praiyot', 'MCA');

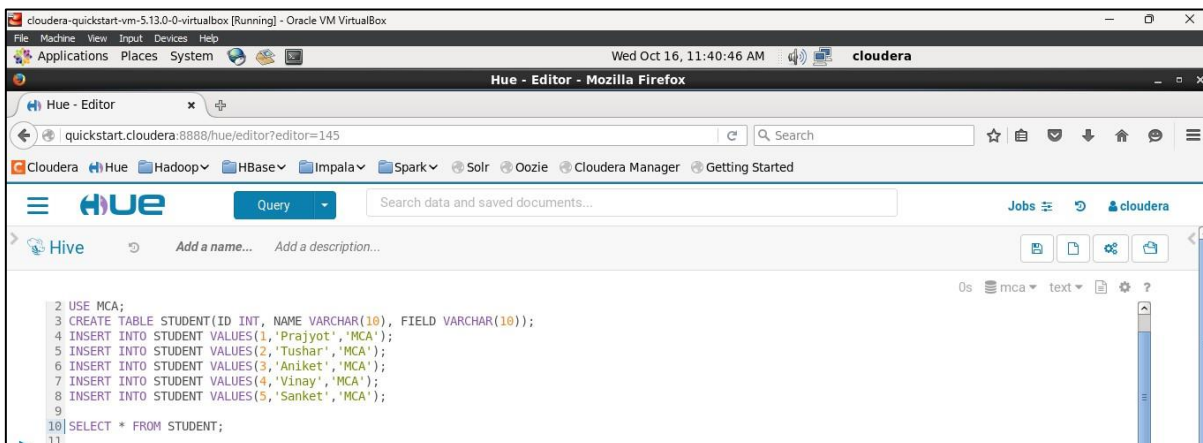
INSERT INTO STUDENT VALUES(2, 'Tushar', 'MCA');

INSERT INTO STUDENT VALUES(3, 'Aniket', 'MCA');

INSERT INTO STUDENT VALUES(4, 'Atharva', 'MCA');

INSERT INTO STUDENT VALUES(5, 'Sanket', 'MCA');

SELECT * FROM STUDENT;



Output:

Query History			Saved Queries			Results (5)		
COLUMNS (3)								
<input checked="" type="checkbox"/>	student.id	INT_TYPE				student.id	student.name	student.field
<input checked="" type="checkbox"/>	student.name	VARCHAR_TYPE				1	1	Prajyot
<input checked="" type="checkbox"/>	student.field	VARCHAR_TYPE				2	2	Tushar
						3	3	Aniket
						4	4	Vinay
						5	5	Sanket

8. Describe the table.

Query History			Saved Queries			Results (3)		
COLUMNS (3)								
<input checked="" type="checkbox"/>	col_name	STRING_TYPE				col_name	data_type	comment
<input checked="" type="checkbox"/>	data_type	STRING_TYPE				1	id	int
<input checked="" type="checkbox"/>	comment	STRING_TYPE				2	name	varchar(10)
						3	field	varchar(10)

9. Describe formatted student.

Query History			Saved Queries			Results (33)		
COLUMNS (3)								
<input checked="" type="checkbox"/>	col_name	STRING_TYPE				col_name	data_type	comment
<input checked="" type="checkbox"/>	data_type	STRING_TYPE				1	# col_name	data_type
<input checked="" type="checkbox"/>	comment	STRING_TYPE				2		NULL
						3	id	int
						4	name	varchar(10)
						5	field	varchar(10)

10. Student order by id desc;

Query History			Saved Queries			Results (5)		
COLUMNS (3)								
<input checked="" type="checkbox"/>	student.id	INT_TYPE				student.id	student.name	student.field
<input checked="" type="checkbox"/>	student.name	VARCHAR_TYPE				1	5	Sanket
<input checked="" type="checkbox"/>	student.field	VARCHAR_TYPE				2	4	Vinay
						3	3	Aniket
						4	2	Tushar
						5	1	Prajyot

11. Select * From student where id=3;

The screenshot shows a SQL query editor with the query `SELECT * FROM STUDENT WHERE ID=3;` and its results. The results table has three columns: `student.id`, `student.name`, and `student.field`. The first row contains the values 3, Aniket, and MCA.

student.id	student.name	student.field
3	Aniket	MCA

12. Select sum(id) from student;

The screenshot shows a SQL query editor with the query `SELECT MAX(ID) FROM STUDENT;` and its results. The results table has one column: `_c0`. The first row contains the value 5.

_c0
5

13. Select Max(id) from student;

The screenshot shows a SQL query editor with the query `SELECT SUM(ID) FROM STUDENT;` and its results. The results table has one column: `_c0`. The first row contains the value 15.

_c0
15

14. Select count(*) from student;

The screenshot shows a SQL query editor with the query `SELECT COUNT(*) FROM STUDENT;` and its results. The results table has one column: `_c0`. The first row contains the value 5.

_c0
5

15. Select name from student limit 3;

19

SELECT NAME FROM STUDENT LIMIT 3;

Query History

Saved Queries

Results (3)

COLUMNS (1)

☐

name

VARCHAR_TYPE

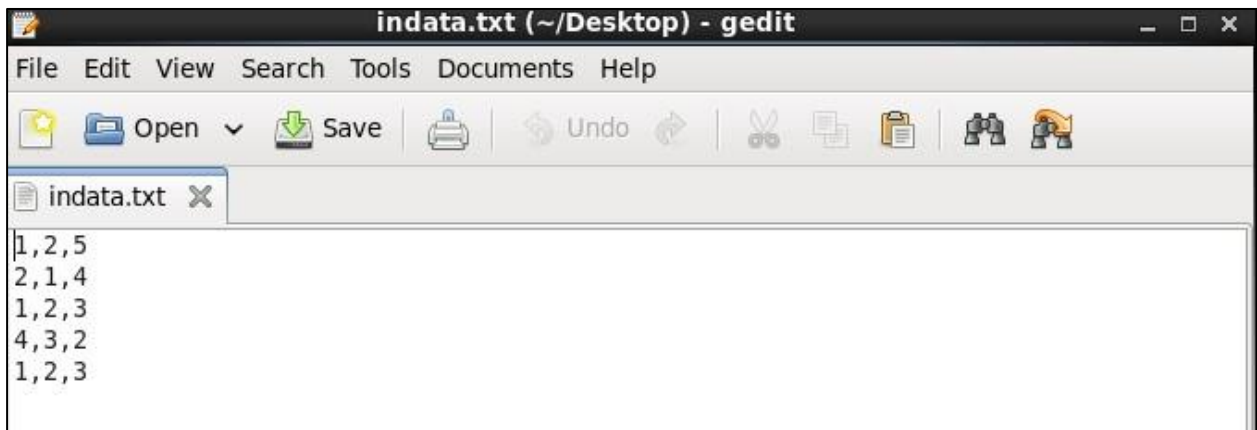
	name
1	Prajyot
2	Tushar
3	Aniket

Practical No : 5

Aim: - PIG List of Commands.

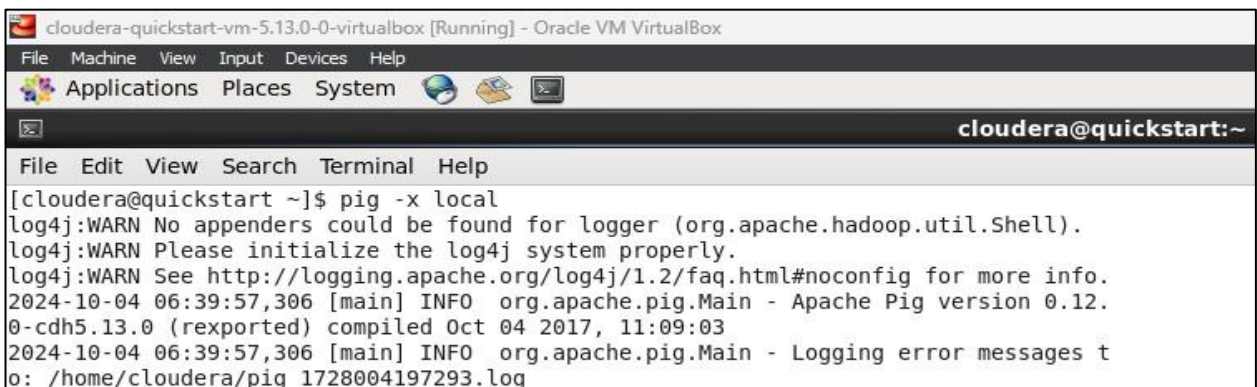
A) Student1.txt

1. Indata.txt



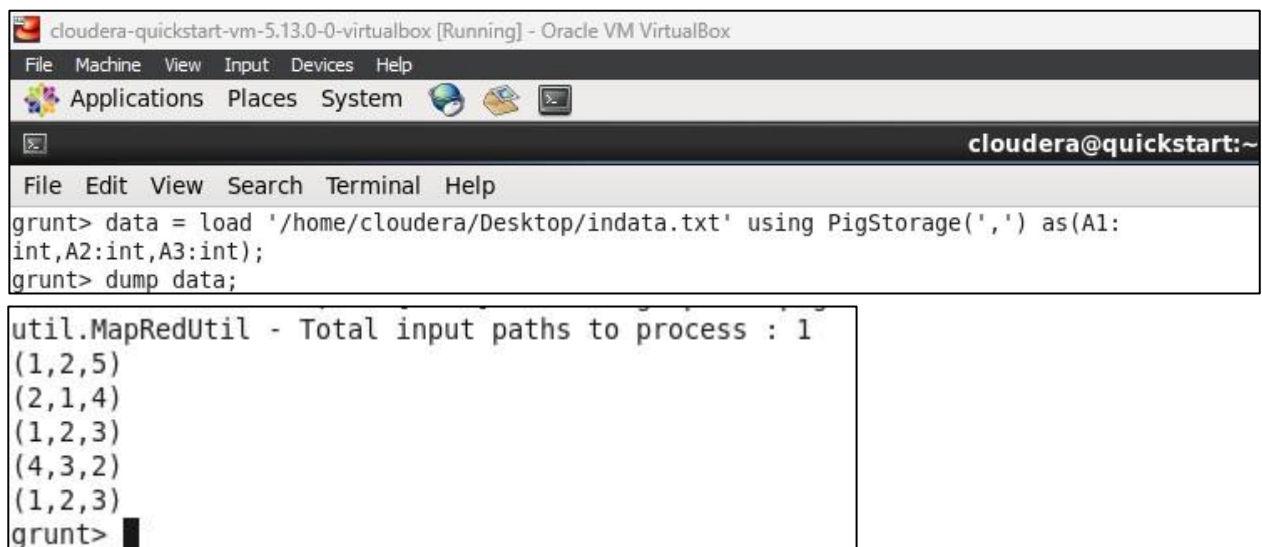
```
indata.txt (~/Desktop) - gedit
File Edit View Search Tools Documents Help
Open Save Undo
indata.txt x
1,2,5
2,1,4
1,2,3
4,3,2
1,2,3
```

2. Pig -x local



```
cloudera-quickstart-vm-5.13.0-0-virtualbox [Running] - Oracle VM VirtualBox
File Machine View Input Devices Help
Applications Places System
cloudera@quickstart:~
File Edit View Search Terminal Help
[cloudera@quickstart ~]$ pig -x local
log4j:WARN No appenders could be found for logger (org.apache.hadoop.util.Shell).
log4j:WARN Please initialize the log4j system properly.
log4j:WARN See http://logging.apache.org/log4j/1.2/faq.html#noconfig for more info.
2024-10-04 06:39:57,306 [main] INFO org.apache.pig.Main - Apache Pig version 0.12.
0-cdh5.13.0 (reexported) compiled Oct 04 2017, 11:09:03
2024-10-04 06:39:57,306 [main] INFO org.apache.pig.Main - Logging error messages t
o: /home/cloudera/pig 1728004197293.log
```

3. Local data and dump data.



```
cloudera-quickstart-vm-5.13.0-0-virtualbox [Running] - Oracle VM VirtualBox
File Machine View Input Devices Help
Applications Places System
cloudera@quickstart:~
File Edit View Search Terminal Help
grunt> data = load '/home/cloudera/Desktop/indata.txt' using PigStorage(',') as(A1:
int,A2:int,A3:int);
grunt> dump data;
util.MapRedUtil - Total input paths to process : 1
(1,2,5)
(2,1,4)
(1,2,3)
(4,3,2)
(1,2,3)
grunt>
```

4. Filter

```
cloudera-quickstart-vm-5.13.0-0-virtualbox [Running] - Oracle VM VirtualBox
File Machine View Input Devices Help
Applications Places System
cloudera@quickstart:~
File Edit View Search Terminal Help
grunt> fitdata = FILTER data by A3==4;
grunt>
grunt> dump fitdata;
```

```
util.MapRedUtil - Total input paths to process : 1
(2,1,4)
grunt>
```

5. Distinct

```
cloudera-quickstart-vm-5.13.0-0-virtualbox [Running] - Oracle VM VirtualBox
File Machine View Input Devices Help
Applications Places System cloudera Fri Oct 4, 7:07:53 AM
cloudera@quickstart:~
File Edit View Search Terminal Help
grunt> result = DISTINCT data;
grunt> dump result;
2024-10-04 07:06:49,303 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pi
g features used in the script: DISTINCT
```

```
util.MapRedUtil - Total input paths to process : 1
(1,2,3)
(1,2,5)
(2,1,4)
(4,3,2)
grunt>
```

6. Foreach

```
cloudera-quickstart-vm-5.13.0-0-virtualbox [Running] - Oracle VM VirtualBox
File Machine View Input Devices Help
Applications Places System
cloudera@quickstart:~
File Edit View Search Terminal Help
grunt> fodata = foreach data generate A1,A2;
grunt>
grunt> dump fodata;
```

```
File Edit View Search Terminal Help
util.MapRedUtil - Total input paths to process : 1
(1,2)
(2,1)
(1,2)
(4,3)
(1,2)
```

7. Load

```
cloudera-quickstart-vm-5.13.0-0-virtualbox [Running] - Oracle VM VirtualBox
File Machine View Input Devices Help
Applications Places System
cloudera
cloudera@quickstart:~
File Edit View Search Terminal Help
grunt> stud = load '/home/cloudera/Desktop/student.txt' using PigStorage(',') as(Name: chararray, Age: int, Grade: chararray);
grunt> dump stud;
```

```
2024-10-04 07:18:02,681 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(Tushar,14, mca)
(Prajyot,53, mca)
(Aniket,42, mca)
(Vinay,10, mca)
(Sanket,53, mca)
(,,)
grunt> █
```

```
student.txt (~/Desktop) - gedit
File Edit View Search Tools Documents Help
Open Save Undo
student.txt
Tushar, 14, mca
Prajyot, 53, mca
Aniket, 42, mca
Vinay, 10, mca
Sanket, 53, mca
```

```

cloudera@quickstart:~
File Edit View Search Terminal Help
grunt> limstud = limit stud 2;
grunt>
grunt> dump limstud;
2022-11-04 02:16:47,068 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig features used in the script: LIMIT
2022-11-04 02:16:47,069 [main] INFO org.apache.pig.newplan.logical.optimizer.LogicalPlanOptimizer - {RULES_ENABLED=[AddForEach, ColumnMapKeyPrune, DuplicateForEachColumnRewrite, GroupByConstParallelSetter, ImplicitSplitInserter, LimitOptimizer, LoadTypeCastInserter, MergeFilter, MergeForEach, NewPartitionFilterOptimizer, PushDownForEachFlatten, PushUpFilter, SplitFilter, StreamTypeCastInserter], RULES_DISABLED=[FilterLogicExpressionSimplifier, PartitionFilterOptimizer]}
2022-11-04 02:16:47,099 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MRCompiler - File concatenation threshold: 100 optimistic? false
2022-11-04 02:16:47,103 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MultiQueryOptimizer - MR plan size before optimization: 2
2022-11-04 02:16:47,105 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MultiQueryOptimizer - MR plan size after optimization: 2

2022-11-04 02:16:47,105 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.Util - Total input paths to process : 1
(Vinay,10, mca)
(Sanket,53, mca)
(,,)
grunt>

```

```

cloudera@quickstart:~
File Edit View Search Terminal Help
grunt>
grunt> grpstud = Group stud by Name;
grunt>
grunt> dump grpstud;
2022-11-04 01:58:34,060 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig features used in the script: GROUP_BY
2022-11-04 01:58:34,061 [main] INFO org.apache.pig.newplan.logical.optimizer.LogicalPlanOptimizer - {RULES_ENABLED=[AddForEach, ColumnMapKeyPrune, DuplicateForEachColumnRewrite, GroupByConstParallelSetter, ImplicitSplitInserter, LimitOptimizer, LoadTypeCastInserter, MergeFilter, MergeForEach, NewPartitionFilterOptimizer, PushDownForEachFlatten, PushUpFilter, SplitFilter, StreamTypeCastInserter], RULES_DISABLED=[FilterLogicExpressionSimplifier, PartitionFilterOptimizer]}

```

```

cloudera@quickstart:~
File Edit View Search Terminal Help
grunt>
grunt> accstud = order stud by Name;
grunt>
grunt> dump accstud;
2022-11-04 02:22:05,189 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig features used in the script: ORDER_BY
2022-11-04 02:22:05,192 [main] INFO org.apache.pig.newplan.logical.optimizer.LogicalPlanOptimizer - {RULES_ENABLED=[AddForEach, ColumnMapKeyPrune, DuplicateForEachColumnRewrite, GroupByConstParallelSetter, ImplicitSplitInserter, LimitOptimizer, LoadTypeCastInserter, MergeFilter, MergeForEach, NewPartitionFilterOptimizer, PushDownForEachFlatten, PushUpFilter, SplitFilter, StreamTypeCastInserter], RULES_DISABLED=[FilterLogicExpressionSimplifier, PartitionFilterOptimizer]}
2022-11-04 02:22:05,208 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MRCompiler - File concatenation threshold: 100 optimistic? false

```

Practical No : 6

AIM : BASIC SPARK COMMANDS

1. Move the data into Hadoop file system

```
[cloudera@quickstart ~]$ hdfs dfs -ls
Found 3 items
drwxr-xr-x   - cloudera cloudera          0 2024-07-26 07:14 DemoDirectory
drwxr-xr-x   - cloudera cloudera          0 2024-09-06 08:58 Demodirector
```

2. Start pyspark in terminal

```
[cloudera@quickstart ~]$ pyspark
Python 2.6.6 (r266:84292, Jul 23 2015, 15:22:56)
[GCC 4.4.7 20120313 (Red Hat 4.4.7-11)] on linux2
Type "help", "copyright", "credits" or "license" for more information.
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel).
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/lib/zookeeper/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/lib/flume-ng/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/lib/parquet/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/lib/avro/avro-tools-1.7.6-cdh5.13.0.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
24/09/21 06:08:08 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
24/09/21 06:08:08 WARN util.Utils: Your hostname, quickstart.cloudera resolves to a loopback address: 127.0.0.1; using 10.0.2.15 instead (on interface eth0)
24/09/21 06:08:08 WARN util.Utils: Set SPARK_LOCAL_IP if you need to bind to another address
Welcome to
```



3. Writing compute contrib function

```
>>> def computeContribs(neighbors ,rank):
...     for neighbor in neighbors : yield(neighbor , rank/len(neighbors))
... 
```

4. Create RDD named links with following command

```
...
>>> links = sc.textFile('web-Google.txt')\
...         .map(lambda line : line.split())\
...         .map(lambda pages : (pages[0],pages[1]))\
...         .distinct()\
...         .groupByKey()\
...         .persist()
```

5. Create ranks rdd storing the ranks data

```
>>> links = sc.textFile('web-Google.txt')\
...         .map(lambda line : line.split())\
...         .map(lambda pages : (pages[0],pages[1]))\
...         .distinct()\
...         .groupByKey()\
...         .persist()
>>> ranks = links.map(lambda (page,neighbors) : (page,1,0))
```

6. Create a loop in order to calculate contribs and ranks

```
>>> for x in xrange(10):
...     contribs = links\
...     .join(ranks)\
...     .flatMap(lambda (page,(neighbors,rank)): computeContribs(neighbors,rank)
... )
...     ranks=contribs\
...     .reduceByKey(lambda v1,v2 :v1+v2)\
...     .map(lambda (page , contrib) : (page,contrib *0.85 +0.15))
```

7. Code to collect all ranks

```
>>> for rank in ranks.collect(): print rank
...
(u'9', 0.26251629106036611)
(u'3', 0.51668019787260355)
(u'6', 0.21345260398897214)
(u'0', 0.52378700337928918)
(u'7', 0.21345260398897214)
(u'1', 0.26962309656705169)
(u'4', 0.49465567868778382)
(u'5', 0.52378700337928918)
(u'8', 0.21345260398897214)
(u'2', 0.26251629106036611)
>>> █
```

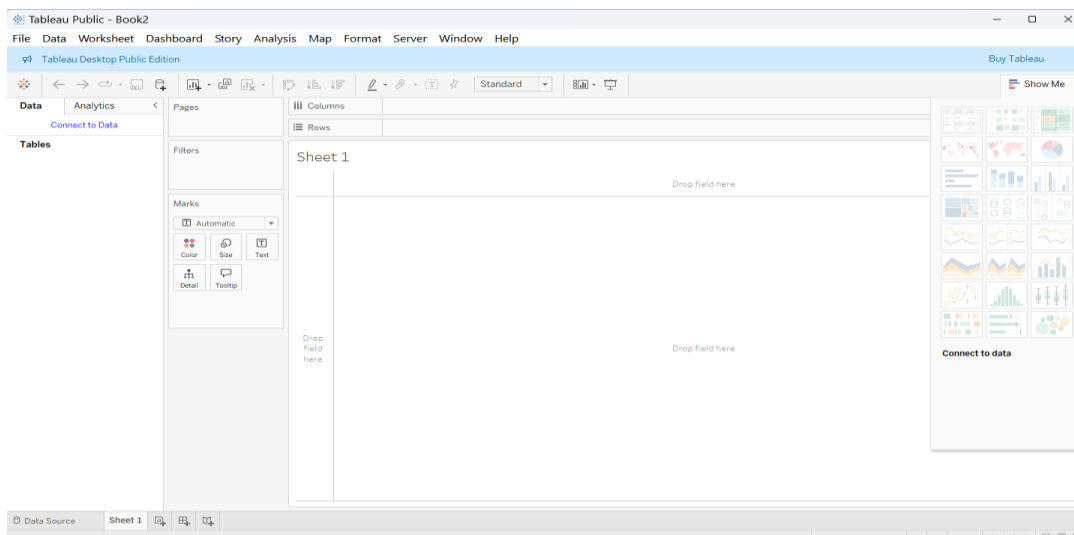
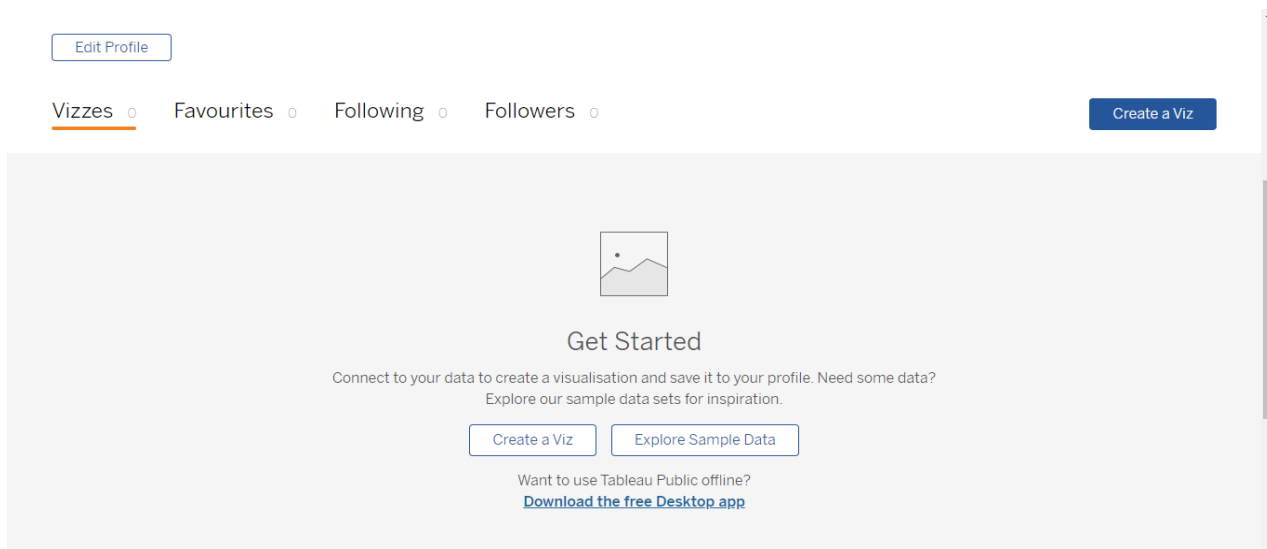

Practical No : 7

AIM : To Demonstrate Visualization Using Tableau

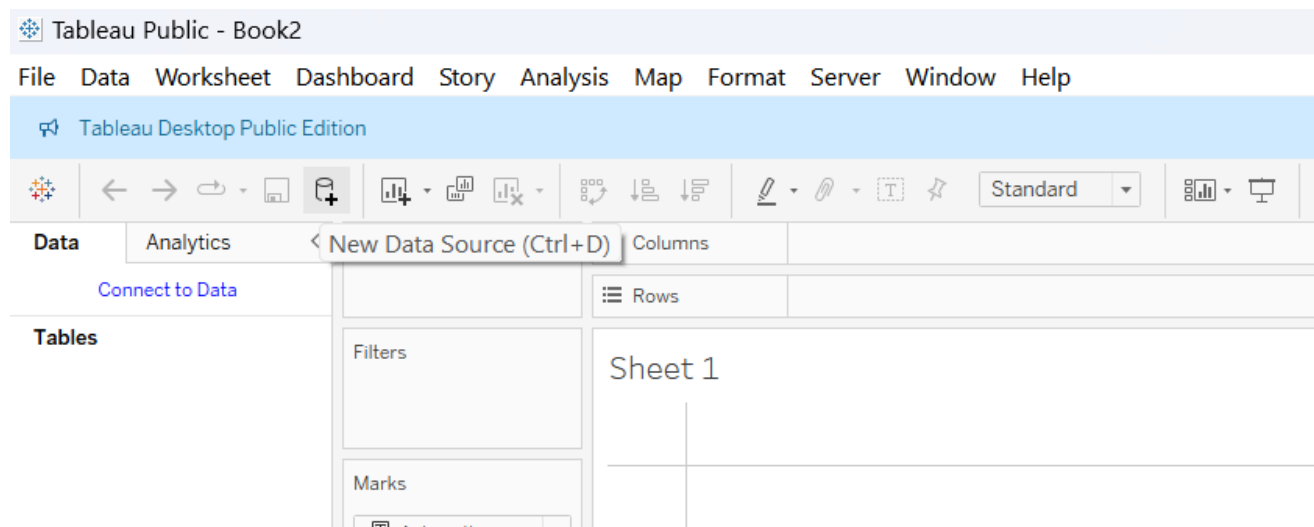
1. Create an Account On Tableau Public



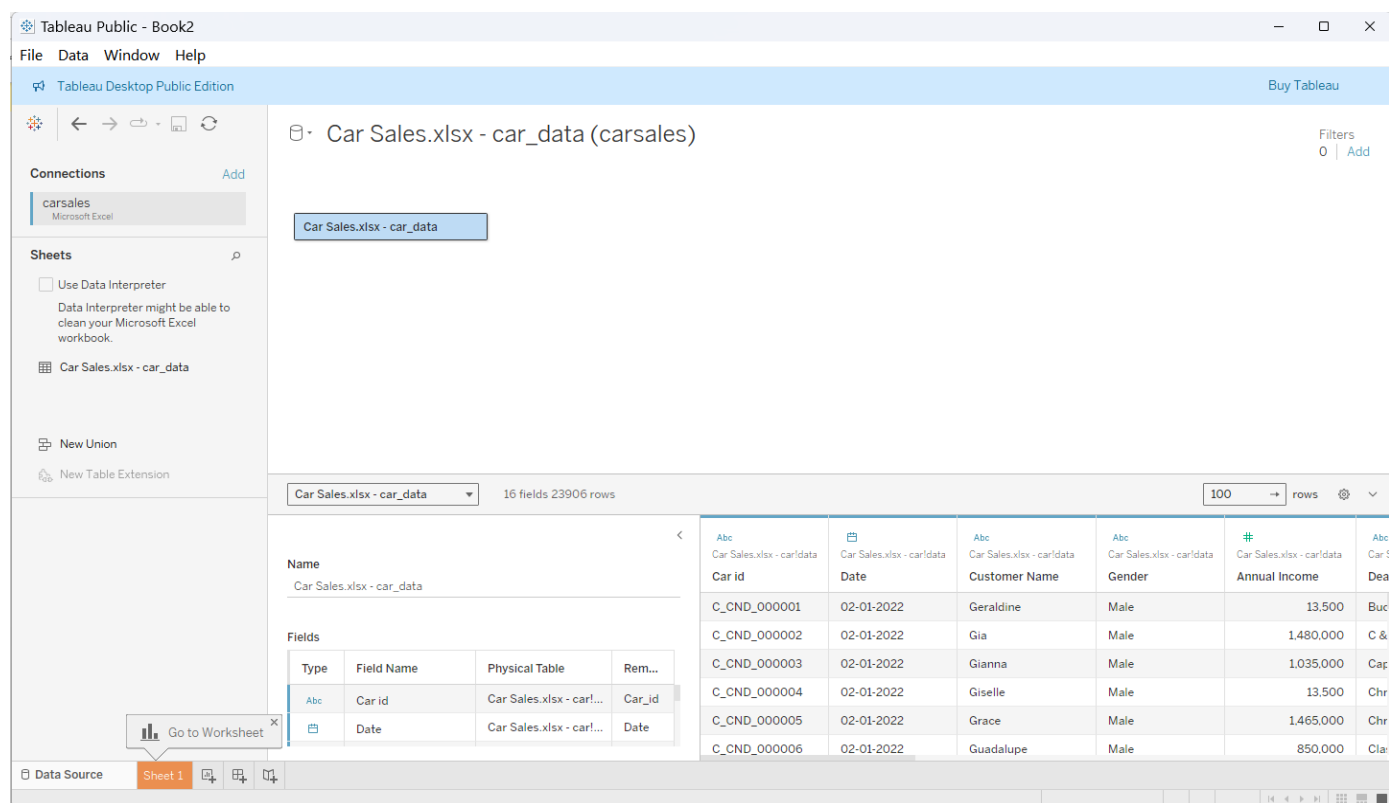
2. Create VIZ



3. Click on the icon Shown below



4. Now Drag and Drop the file you want to use



5. Drag and Drop any field you want to use on row

Tableau Desktop Public Edition

Columns: Customer Name

Rows: (Empty)

Marks: Automatic

Tables:

- Body Style
- Car id
- Color
- Company
- Customer Name
- Date
- Dealer Name
- Dealer No
- Dealer Region
- Engine
- Gender
- Model
- Transmission
- Measure Names
- Annual Income
- Phone
- Price (\$)
- Car Sales.xlsx - car_data (...)
- Measure Values

6. Drag and Drop csv file onto data part

- # Annual Income
- # Phone
- # Price (\$)
- # Car Sales.xlsx - car_data (...)
- # Measure Values

Tableau Desktop Public Edition

Columns: Dealer Region

Rows: Customer Name

Marks: Automatic

Tables:

- Body Style
- Car id
- Color
- Company
- Customer Name
- Date
- Dealer Name
- Dealer No
- Dealer Region
- Engine
- Gender
- Model
- Transmission
- Measure Names
- Annual Income
- Phone
- Price (\$)
- Car Sales.xlsx - car_data (...)
- Measure Values

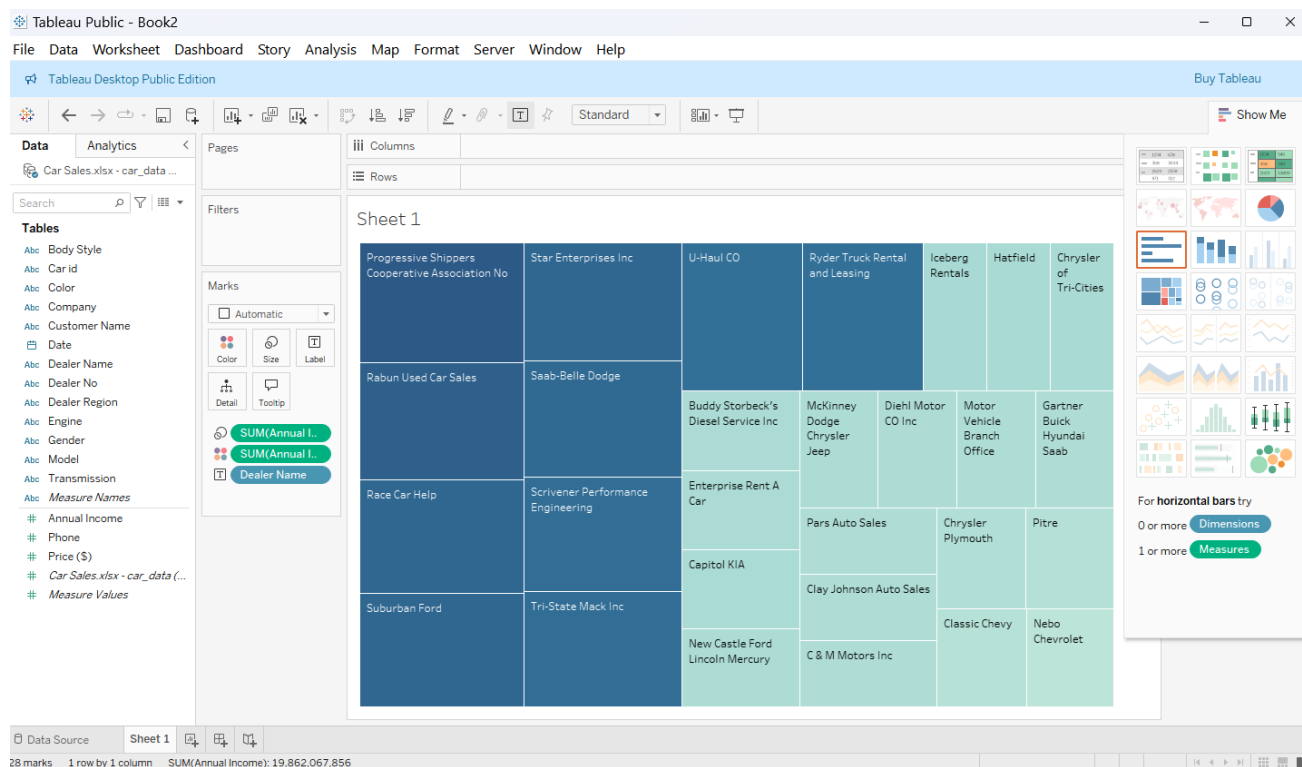
Sheet 1

Customer Name	Aurora	Austin	Greenvi..	Janesvi..	Middlet..	Pasco	Scottsd..
Aahil						2	
Aaliyah	3	2	2	2	2	4	2
Aarav	2	1		2		1	2
Aaron	4	8	7	4	7	4	5
Aarya			1	1			
Aayan	1				1	1	2
Abby		2	2	3	1	2	1
Abdal		1					
Abdelatif		1					
Abderramine		1					
Abdiel		2	1	1	1	1	1
Abdjadi						1	
Abdoul	3		2			3	
Abdoulaye		1	3	4		2	
Abdul	1	2		1	1	1	2
Abdullah	1	3	2	2	1	1	1
Abel	3	2		1	3	2	1
Abigail	3	7	3	6	3	5	9
Abraham	3	2	3	4	5	4	2
Abrielle			1				1
Abriil							
Abthu-Mustapha	1					1	
Achille				1			1
Ada	2	4		4	2		1

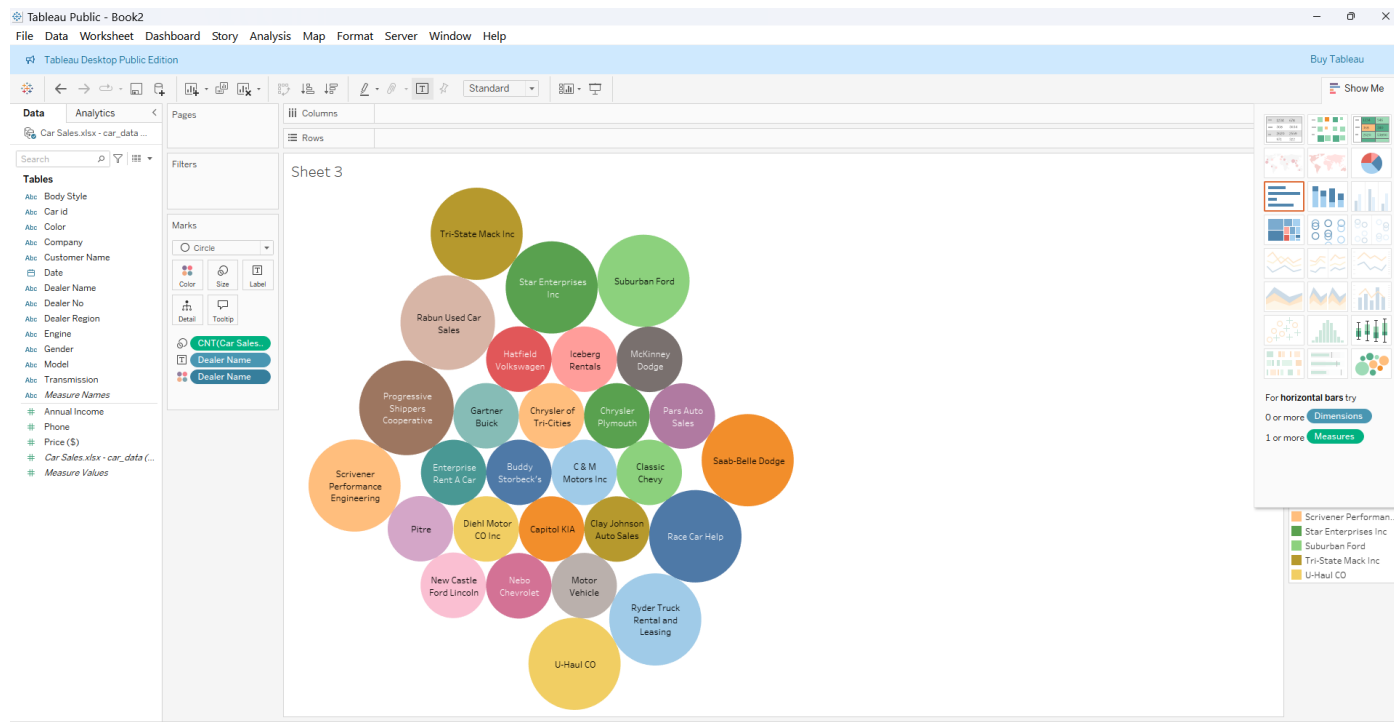
10070 marks 3022 rows by 7 columns SUM of CNT(Car Sales.xlsx - car_data): 23.906

7. Analysing Using Charts

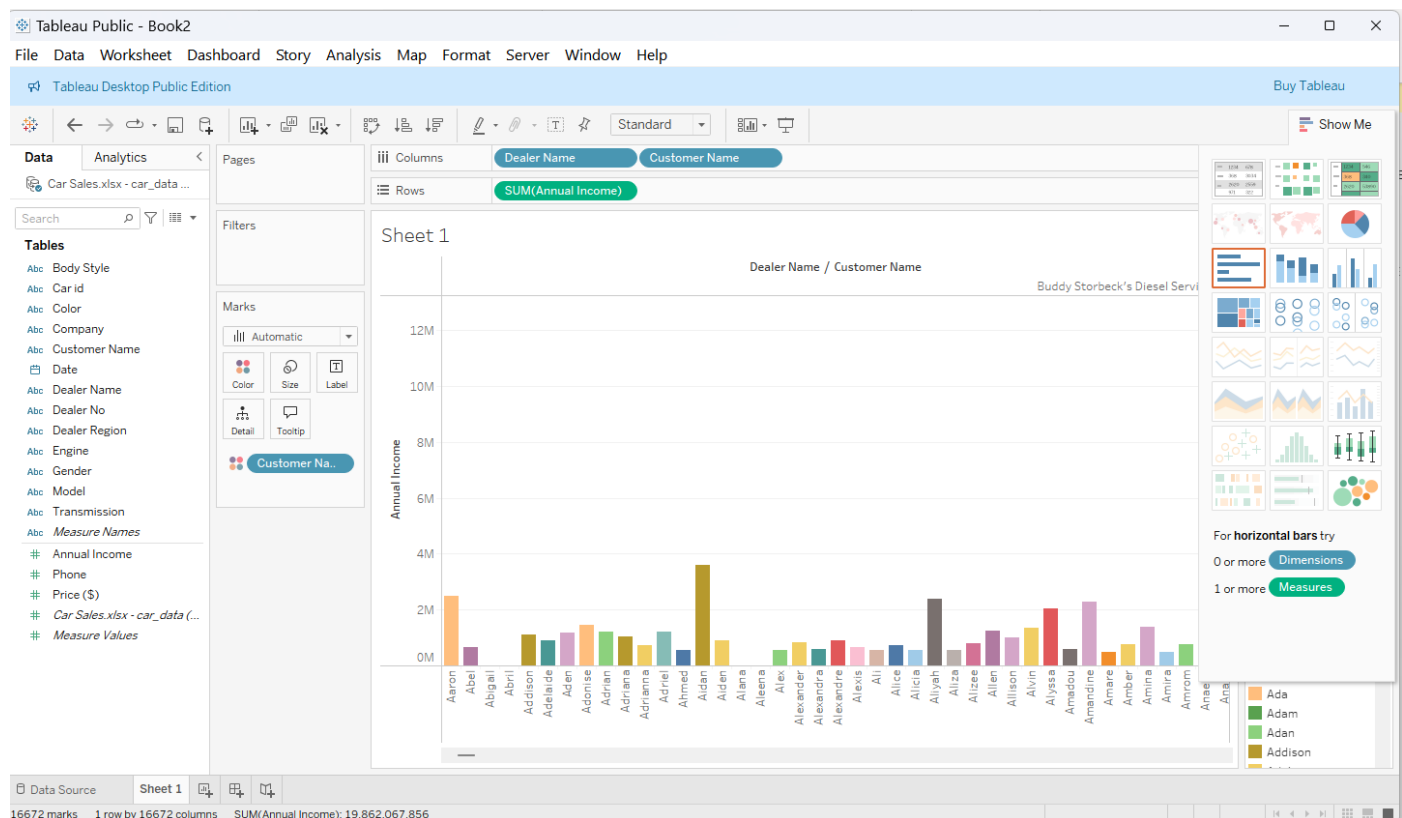
◆ TREEMAP



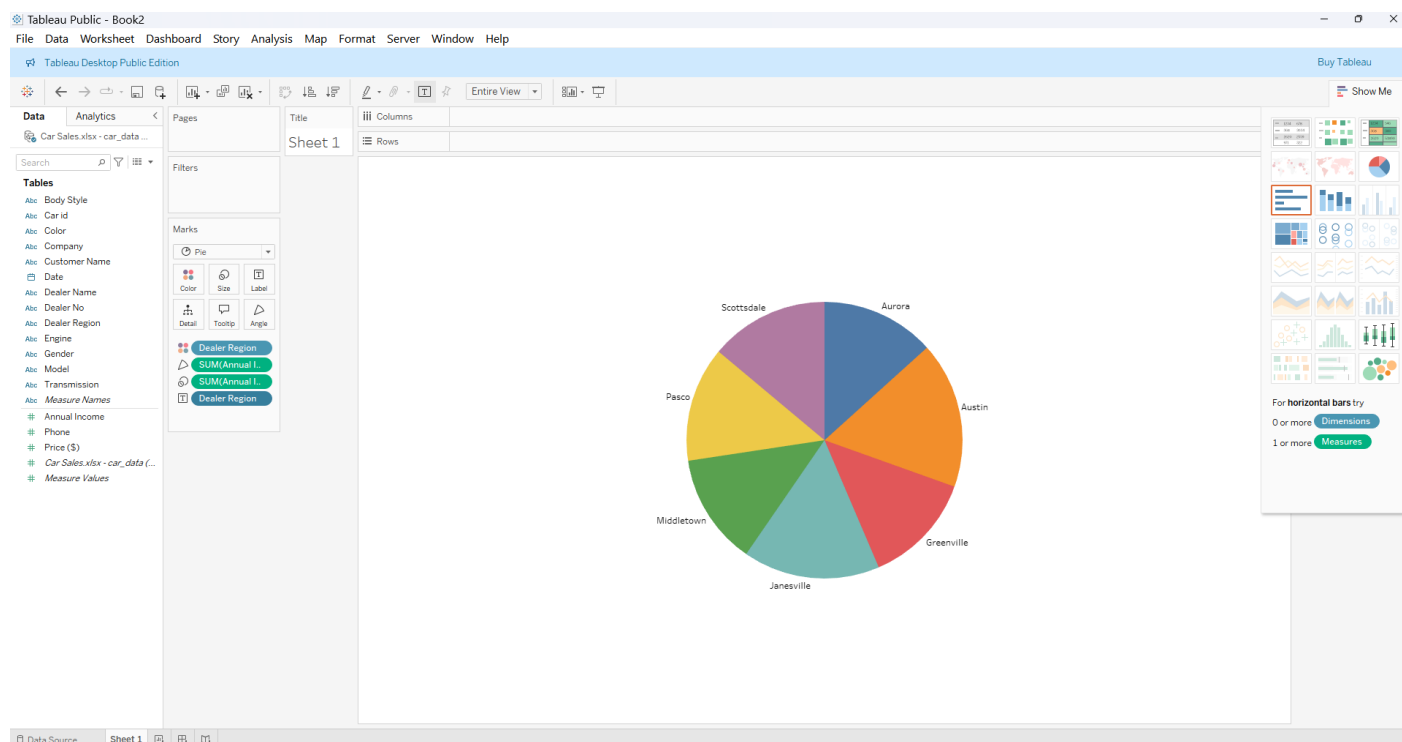
◆ PACKED BUBBLES



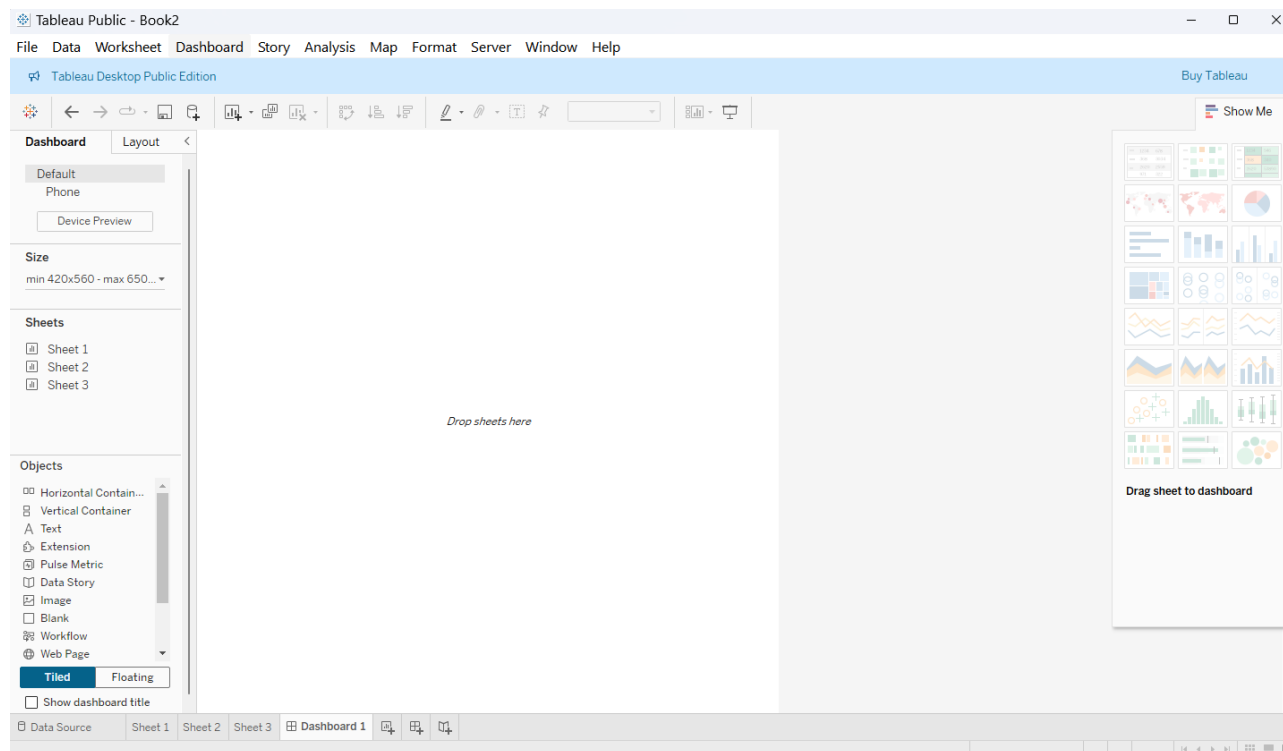
◆ SIDE-BY-SIDE BARS



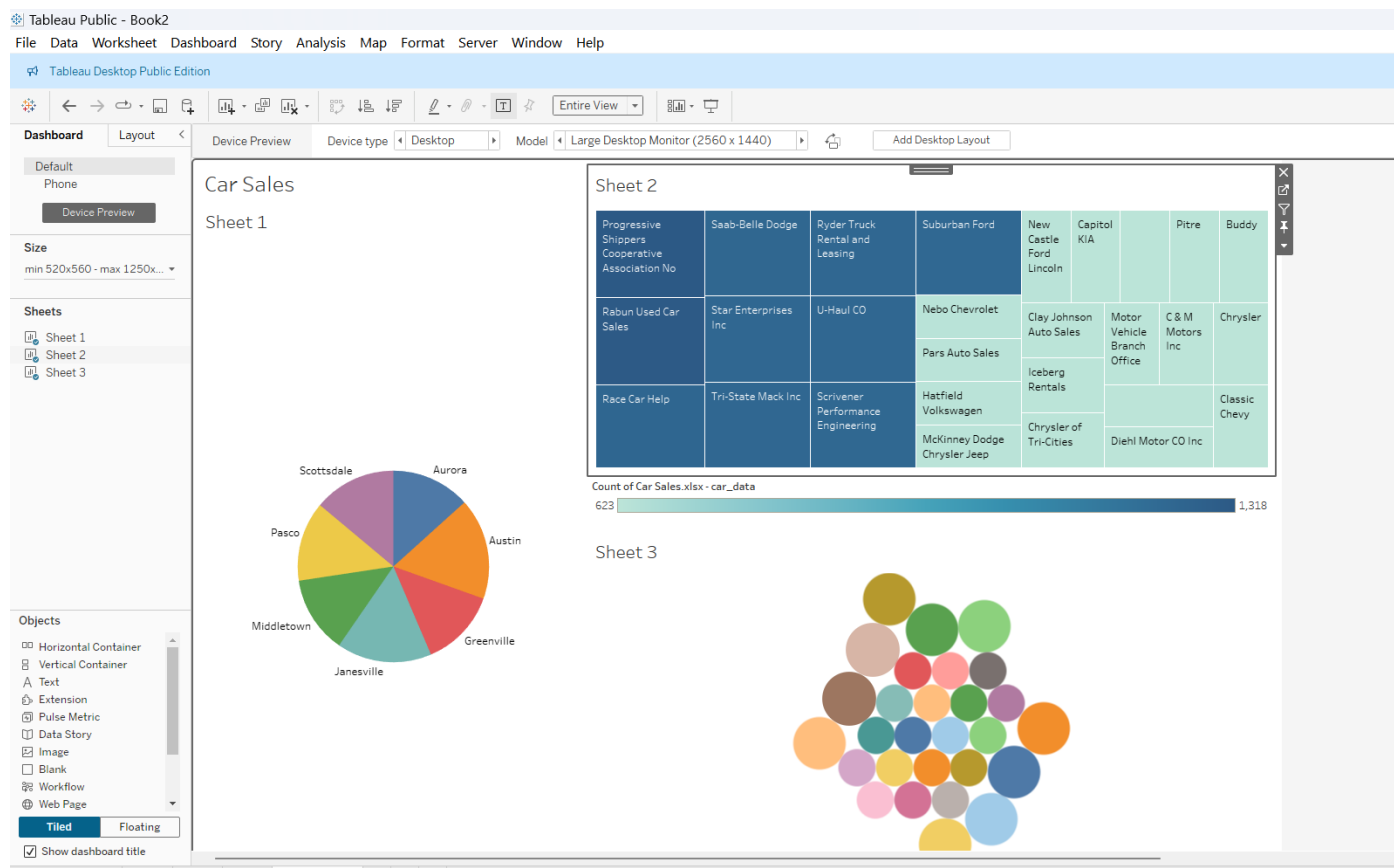
◆ PIE CHARTS



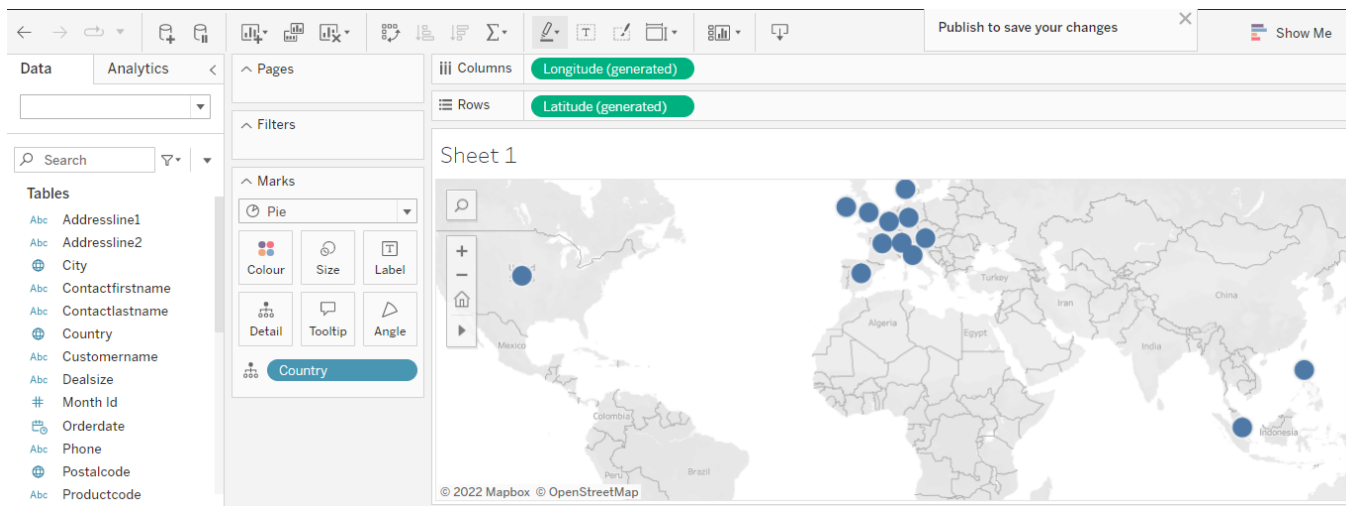
8. Creating Dashboard



● Drag and drop your sheet at your left side at main working space

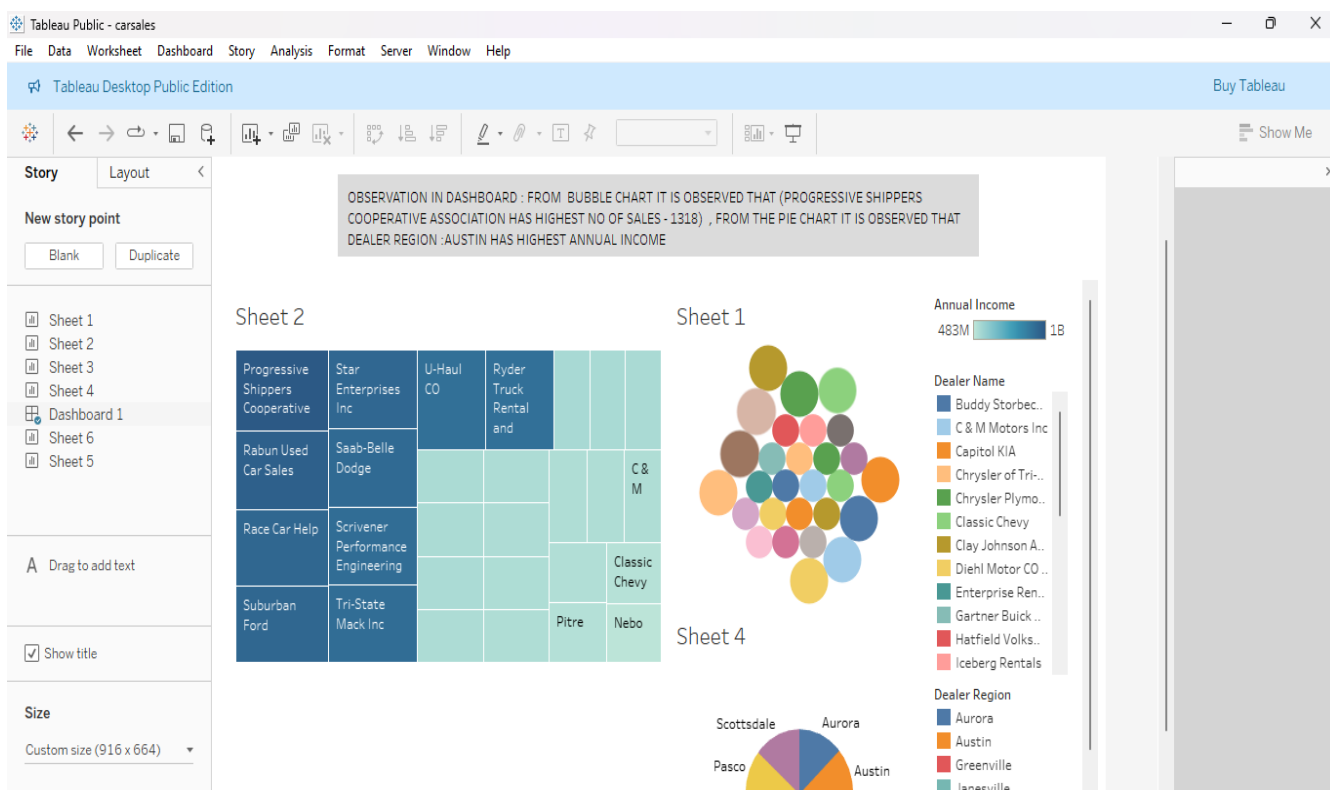


9. Working with Maps



10. Telling Stories With Tableau

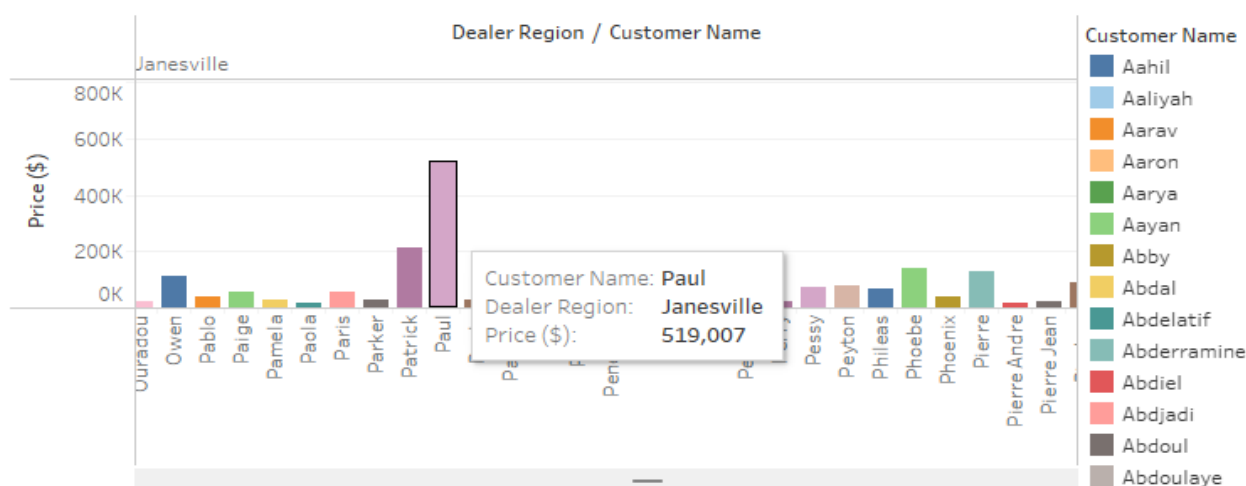
•ANALYSIS OF DASHBOARD



• CUSTOMER NAME PAID HIGHEST FOR CARS

Story 2


FROM SIDE BY SIDE BARS IT IS OBSERVED THAT CUSTOMER NAME : PAUL HAS HIGHEST PAID PRICE FOR CARS WHICH IS IN THE DEALER REGION JANESVILLE



• DEALER REGION HAVING HIGHEST ANNUAL INCOME

Story 3

OBSERVATION FROM PIE CHART IS DEALER REGION AUSTIN HAS HIGHEST ANNUAL INCOME AS COMPARE TO OTHER DEALER REGION. DEALER REGION JANESVILLE HAS SECOND HIGHEST ANNUAL INCOME



- **MAP OBSERVATION HAVING DEALER REGION AND ANNUAL INCOME**

Story 4

BELOW MAPS SHOWS THE DEALER REGION AND ANNUAL INCOME

